

CMPT 413

Computational Linguistics

Anoop Sarkar

<http://www.cs.sfu.ca/~anoop>


Sequence Learning

- British Left Waffles on Falkland Islands
 - (N, N, V, P, N, N)
 - (N, V, N, P, N, N)
- Segmentation 中国十四个边境开放城市经济建设成就显著
 - (s, b, i, b, i, b, b, i, b, i, b, i, b, i, b, i, b, i)

中国 十 四 个 边 境 开 放 城 市 经 济 建 设 成 就 显 著

China 's 14 open border cities marked economic achievements

Sequence Learning




3 states: N, V, P

Observation sequence: (o_1, \dots, o_6)


State sequence (6+1): $(Start, N, N, V, P, N, N)$

Finite State Machines



Mealy Machine

Finite State Machines

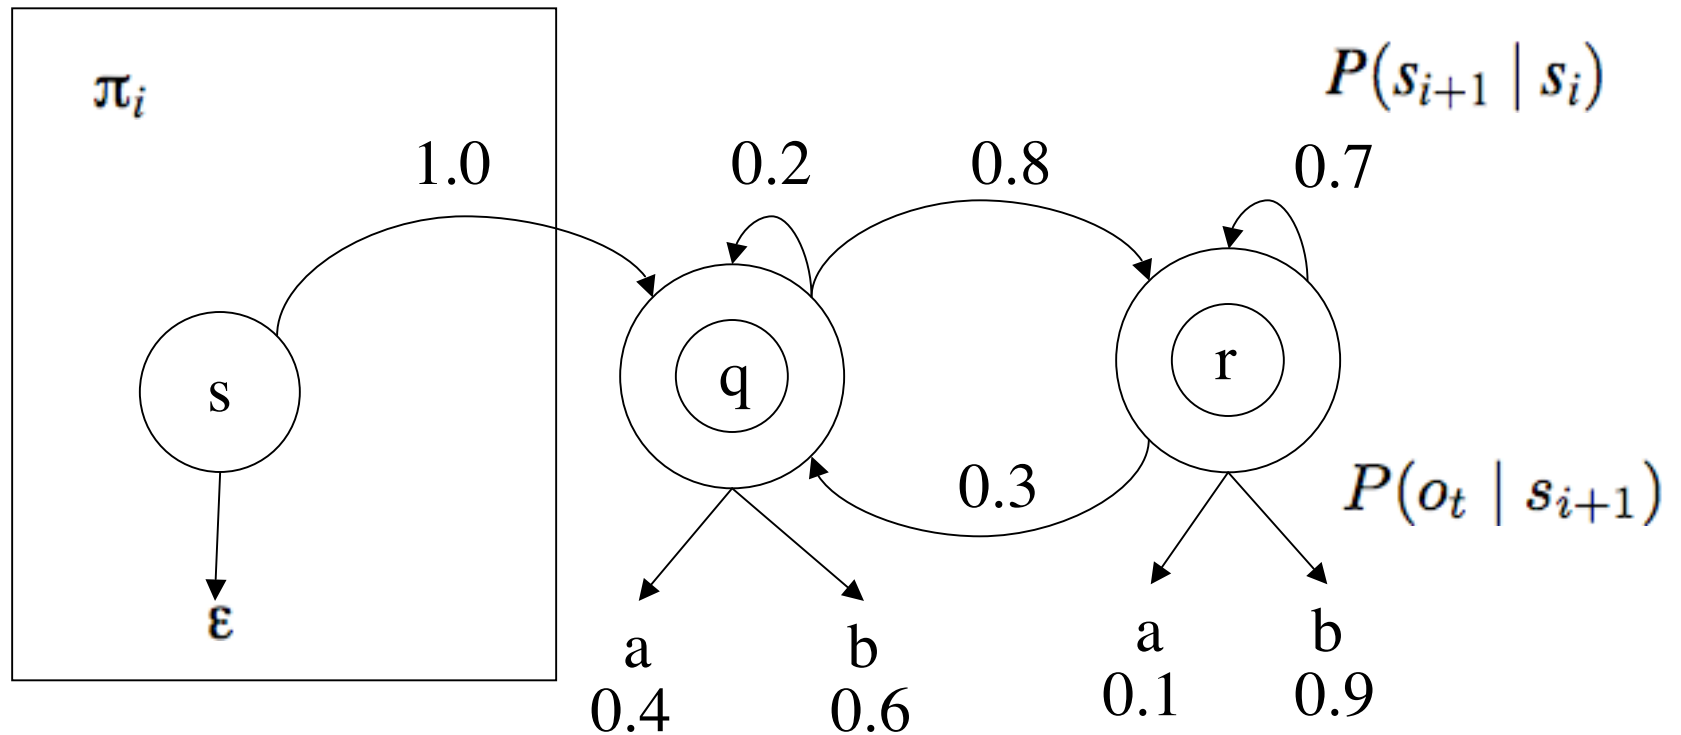


Moore Machine

Probabilistic FSMs

- Each transition is associated with a *transition probability*
- Each emission is associated with an *emission probability*
- Two conditions:
 - All outgoing transition arcs from a state must sum to 1
 - All emission arcs from a state must sum to 1

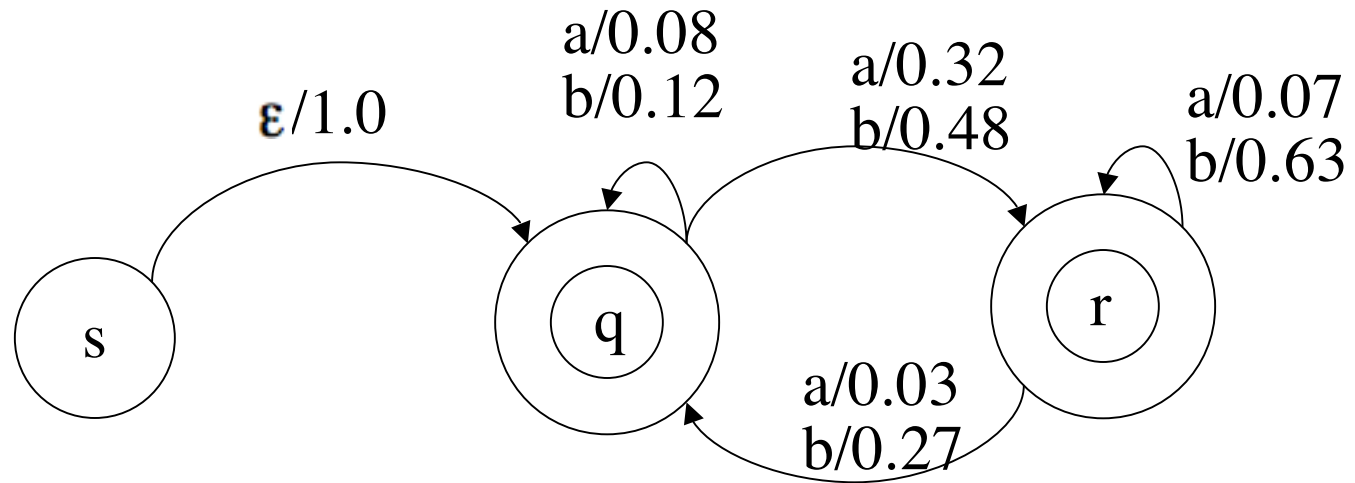
Probabilistic FSMs



$$\sum_x P(q \rightarrow x) = P(q \rightarrow r) + P(q \rightarrow q) = 1.0$$

$$\sum_x P(\text{emit}(q, x)) = P(\text{emit}(q, a)) + P(\text{emit}(q, b)) = 1.0$$

Probabilistic FSMs




Hidden Markov Models

- There are n states $s_1, \dots, s_i, \dots, s_n$
- The emissions are observed (input data)
- Observation sequence $\mathbf{O}=(o_1, \dots, o_t, \dots, o_T)$
- The states are not directly observed (hidden)
- Data does not directly tell us which state X_t is linked with observation o_t

$$X_t \in \{s_1, \dots, s_n\}$$

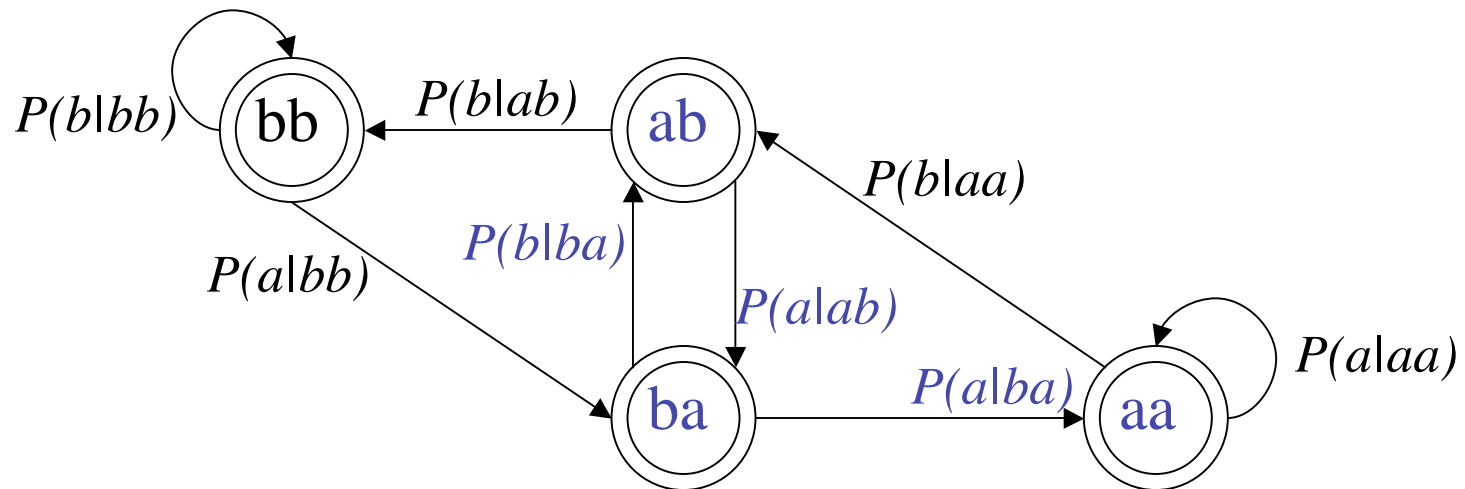
Markov Chains vs. HMMs

- For observation sequence $babaa$
i.e.: $o_1=b, o_2=a, \dots, o_5=a$
- Compute $P(babaa)$ using a bigram model
 $P(b)*P(a|b)*P(b|a)*P(a|b)*P(a|a)$
- Equivalent Markov chain:




Markov Chains vs. HMMs

- For observation sequence $babaa$
i.e.: $o_1=b, o_2=a, \dots, o_5=a$
- Compute $P(babaa)$ using a trigram model
 $P(ba)*P(b|ba)*P(a|ab)*P(a|ba)$
- Equivalent Markov chain:



Markov Chains vs. HMMs

- For observation sequence $babaa$
i.e.: $o_1=b, o_2=a, \dots, o_5=a$
- Compute $P(babaa)$ using a trigram model
 $P(ba)*P(b|ba)*P(alab)*P(alba)$
- Equivalent Markov chain:



Markov Chains vs. HMMs

- Given an observation sequence

$$\mathbf{O}=(o_1, \dots, o_t, \dots, o_T)$$

- An n th order Markov Chain or n -gram model computes the probability

$$P(o_1, \dots, o_t, \dots, o_T)$$

- An HMM computes the probability

$$P(X_1, \dots, X_{T+1}, o_1, \dots, o_T) \text{ where the state sequence is } \textit{hidden}$$

Properties of HMMs

- Markov assumption

$$P(X_t = s_i \mid \dots, X_{t-1} = s_j)$$

- Stationary distribution

$$P(X_t = s_i \mid X_{t-1} = s_j) = P(X_{t+l} = s_i \mid X_{t+l-1} = s_j)$$

HMM Algorithms

- HMM as language model: compute probability of given observation sequence
- HMM as parser: compute the best sequence of states for a given observation sequence
- HMM as learner: given a set of observation sequences, learn its distribution, i.e. learn the transition and emission probabilities

HMM Algorithms

- HMM as language model: compute probability of given observation sequence
- Compute $P(o_1, \dots, o_T)$ from the probability $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$

$$= \prod_{t=1}^T P(X_{t+1} = s_j \mid X_t = s_i) \times P(o_t = k \mid X_{t+1} = s_j)$$

$$P(o_1, \dots, o_T) = \sum_{X_1, \dots, X_{T+1}} P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$$


HMM Algorithms

- HMM as parser: compute the best sequence of states for a given observation sequence
- Compute best path X_1, \dots, X_{T+1} from the probability $P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$

Best state sequence X_1^*, \dots, X_{T+1}^*

$$= \operatorname{argmax}_{X_1, \dots, X_{T+1}} P(X_1, \dots, X_{T+1}, o_1, \dots, o_T)$$

Best Path (Viterbi) Algorithm



- Key Idea 1: storing just the best path doesn't work
- Key Idea 2: store the best path upto *each* state

Forward-Backward Algorithm

- Algorithm that finds the transition and emission probabilities using training data that *does not have* hidden states provided
- Set the probabilities (for all parameters in the HMM) so that the training data T is assigned highest $P(T)$ value (or lowest $H(T)$, entropy value)
- This is called the maximum likelihood value over all possible hidden state sequences for the training data
- Exploits the fact that some transitions and resulting observations will occur more frequently than others in the training data

Forward-backward Algorithm

- Consider input $o_1, \dots, o_t, \dots, o_T$ where each o_t is from a set of symbols $V = \{1, \dots, k, \dots, K\}$
- Let π_i be the probability of state i being a start state (for simplicity, π_i is not discussed further)
- Let $a_{i,j}$ be the transition probability:
$$P(X_{t+1} = s_j \mid X_t = s_i) \quad |S|^2 \text{ distinct } a_{i,j} \text{ values}$$
- Let $b_{j,k}$ be the emission probability:
$$P(o_t = k \mid X_{t+1} = s_j) \quad |S| \times |V| \text{ distinct } b_{j,k} \text{ values}$$
- Probability of going from state s_i to state s_j while observing input o_t is simply $a_{i,j} \times b_{j,k}$

Forward-backward Algorithm

- The algorithm starts with an initial setting for the probabilities in a and b
- We are provided with training data which consists of observation sequence(s): $o_1, \dots, o_t, \dots, o_T$
- The probability $P(o_1, \dots, o_T)$ depends on the values in a and b
- For given observation sequence(s), different transitions/emissions will be visited with different frequencies

Forward-backward Algorithm

- For every path through the HMM, we count how many transitions occurred from state i to state j on observation o_t
- Then (loosely speaking) we reward those transitions (and emissions) which have high *expected* frequency and penalize the competing transitions
- Expected frequency means we multiply the frequency with the current probability (taken from a and b)


Forward-backward Algorithm

- $P(o_1, \dots, o_T)$ is the expected frequency of visiting all transitions and so the new frequency is the expected occurrence of a transition divided by $P(o_1, \dots, o_T)$
- This gives us new values for all probabilities: a' and b' and we set a and b to these new values
- Compute $P(o_1, \dots, o_T)$. If the value is unchanged from before iteration then stop (convergence)
- Otherwise iterate (the entire procedure) with new values for a and b

Forward-backward Algorithm

- How to compute expected frequency over all paths efficiently (*reuse dynamic programming idea from Viterbi algorithm*)
- For input $o_1, \dots, o_t, \dots, o_T$ where $o_t \in V = \{1, \dots, k, \dots, K\}$
- For every path from a start state to state i we can compute the probability of observing o_1, \dots, o_{t-1}
- Let $\alpha_i(t)$ be the sum of all these probabilities
- For every path from state j to a final state we can compute the probability of observing o_{t+1}, \dots, o_T
- Let $\beta_j(t+1)$ be the sum of all these probabilities

Forward-Backward Algorithm




$$\alpha_i(t) = \sum_{k=1}^N a_{k,i} \cdot b_{i,o_{t-1}} \cdot \alpha_k(t-1)$$

$$\beta_j(t+1) = \sum_{m=1}^N a_{j,m} \cdot b_{m,o_{t+1}} \cdot \beta_m(t+2)$$

$$P(o_1, \dots, o_T) = \sum_{i=1}^N \alpha_i(T+1) = \sum_{i=1}^N \pi_i \cdot \beta_i(1)$$

Forward-Backward Algorithm



$$\hat{f}(i, j, o_t) = \frac{\alpha_i(t) \cdot a_{i,j} \cdot b_{j,o_t} \cdot \beta_j(t+1)}{P(o_1, \dots, o_T)} \quad \hat{f}(i, j) = \sum_{t=1}^T \hat{f}(i, j, o_t)$$

$$a'_{i,j} = \frac{\hat{f}(i, j)}{\sum_{j=1}^N \hat{f}(i, j)} \quad b'_{j,k} = \frac{\sum_{i=1}^N \sum_{t=1}^T \hat{f}(i, j, o_t = k)}{\sum_{i=1}^N \sum_{j=1}^N \hat{f}(i, j)}$$

Forward-Backward Algorithm

- Each iteration provides new values for all the *parameters*
- But are the new parameters any better? How can we tell?
- Compute probability of the training data
- For HMMs, Baum 1977 shows that the probability will always be non-decreasing (later generalized to the more general EM algorithm)
- Same as cross-entropy is non-increasing

$$KL(\mu_{i+1} || D) \leq KL(\mu_i || D)$$