

# CMPT-413

## Computational Linguistics

Anoop Sarkar

<http://www.cs.sfu.ca/~anoop>

# Human Supervision in Part of Speech Tagging

- In unseen data, we wish to find the part of speech tags:

Input: *In 1994 , Hartnett said*

Output: *In\_IN 1994\_CD ,\_, Hartnett\_NNP said\_VBD*

- The set of part of speech tags are decided by experts
- The experts also have to provide adequate amounts of data in which the part of speech tags have been listed for each word in context.
- This general approach is called **supervised learning** since the training data is provided by humans.

# Trigram Models for Part of Speech Tagging

In\_IN 1994\_CD ,\_, Hartnett\_NNP said\_VBD

THE\_DT BONEYARD\_NNP Northrop\_NNP Grumman\_NNP 's\_POS modest\_JJ  
flight\_NN museum\_NN occupies\_VBZ a\_DT corner\_NN of\_IN one\_CD of\_IN  
its\_PRP\\$ power-seat\_NN adjusters\_NNS ,\_, door\_NN trim\_JJ now\_RB  
made\_VBN in\_IN South\_NNP Korea\_NNP 's\_POS antiquated\_JJ coal-fired\_JJ  
power\_NN plant\_NN in\_IN Canada\_NNP ,\_, to\_TO a\_DT 11.9\_CD million\_CD  
mark\_NN investment\_NN in\_IN Samsung\_NNP 's\_POS Sachon\_NNP plant\_NN  
in\_IN Taiwan\_NNP as\_IN part\_NN of\_IN a\_DT steam\_NN turbine\_NN ,\_,  
a\_DT new\_JJ high-yielding\_JJ rice\_NN plant\_NN was\_VBD reorganized\_VBN  
into\_IN a\_DT big\_JJ expansion\_NN of\_IN a\_DT fuel-fabrication\_NN  
plant\_NN near\_IN Nagoya\_NNP in\_IN Aichi\_NNP Prefecture\_NNP

From\_IN October\_NNP ,\_, when\_WRB they\_PRP  
do\_VBP not\_RB need\_VB it\_PRP

Borges gives a vague reference to some work by Franz Kuhn allegedly commenting on the classification of animals by a Chinese encyclopedia called the \_Celestial Emporium of Benevolent Knowledge\_.

>> ... animals are divided into:

- (a) those that belong to the Emperor,
  - (b) embalmed ones,
  - (c) those that are trained,
  - (d) suckling pigs,
  - (e) mermaids,
  - (f) fabulous ones,
  - (g) stray dogs,
  - (h) those that are included in this classification,
  - (i) those that tremble as if they were mad,
  - (j) innumerable ones,
  - (k) those drawn with a very fine camel brush,
  - (l) others,
  - (m) those that have just broken a flower vase,
  - (n) those that resemble flies from a distance. <<
- Jorge Luis Borges, "Other Inquisitions"

## Part of Speech Tagging using Trigram Models

- Let the input sentence (word sequence) be  $w_0, w_1, \dots, w_n$
- Let the most likely tag sequence be  $T^* = t_0^*, t_1^*, \dots, t_n^*$
- In order to compare all possible tag sequences we build a probability model:

$$P(t_0, t_1, \dots, t_n \mid w_0, w_1, \dots, w_n)$$

# Part of Speech Tagging using Trigram Models

- The best (or most likely) tag sequence is:

$$T^* = \arg \max_{t_0, \dots, t_n} P(t_0, \dots, t_n \mid w_0, \dots, w_n)$$

$$P(t_0, \dots, t_n \mid w_0, \dots, w_n)$$

$$= \frac{P(w_0, \dots, w_n \mid t_0, \dots, t_n) \times P(t_0, \dots, t_n)}{P(w_0, \dots, w_n)} \text{(Bayes Rule)}$$

$$= P(w_0, \dots, w_n \mid t_0, \dots, t_n) \times P(t_0, \dots, t_n)$$

## Part of Speech Tagging using Trigram Models

$$\begin{aligned} P(w_0, \dots, w_n \mid t_0, \dots, t_n) \\ &= P(w_0 \mid t_0) \times P(w_1 \mid t_1) \times \dots \times P(w_n \mid t_n) \\ &= \prod_{i=0}^n P(w_i \mid t_i) \end{aligned}$$

$$\begin{aligned} P(t_0, \dots, t_n) \\ &= P(t_0) \times P(t_1 \mid t_0) \times P(t_2 \mid t_0, t_1) \times \dots \times P(t_n \mid t_{n-2}, t_{n-1}) \\ &= P(t_0) \times P(t_1 \mid t_0) \times \prod_{i=2}^n P(t_i \mid t_{i-2}, t_{i-1}) \end{aligned}$$

## Part of Speech Tagging using Trigram Models

$$P(t_0, \dots, t_n \mid w_0, \dots, w_n)$$

$$= P(w_0, \dots, w_n \mid t_0, \dots, t_n) \times P(t_0, \dots, t_n)$$

$$= \left( \prod_{i=0}^n P(w_i \mid t_i) \right) \times \left( P(t_0) \times P(t_1 \mid t_0) \times \prod_{i=2}^n P(t_i \mid t_{i-2}, t_{i-1}) \right)$$

$$= \prod_{i=0}^n P(w_i \mid t_i) \times P(t_i \mid t_{i-2}, t_{i-1})$$



## Part of Speech Tagging using Trigram Models

$$P(t_0, \dots, t_n \mid w_0, \dots, w_n) = \prod_{i=0}^n P(w_i \mid t_i) \times P(t_i \mid t_{i-2}, t_{i-1})$$

- This allows us to represent tagging as a Hidden Markov Model (*hmm*).
- The advantage is that we can reuse efficient *hmm* algorithms like Viterbi to find the most likely tag sequence for a given word sequence.
- However, instead of using Forward-Backward to find the values of  $P(w_i \mid t_i)$  and  $P(t_i \mid t_{i-2}, t_{i-1})$  we directly use frequencies from human labelled training data

## Part of Speech Tagging using Trigram Models

$$P(t_0, \dots, t_n \mid w_0, \dots, w_n) = \prod_{i=0}^n P(w_i \mid t_i) \times P(t_i \mid t_{i-2}, t_{i-1})$$

- We can construct a *hmm* that is equivalent to the above model. Exactly the same construction as equivalence of Markov chains with *n*-gram models.
  - Except instead of pairs of words we have pairs of tags as states in the Markov chain.
  - And we add the emission probability to each state to extend the Markov chain to a *hmm*.

## Part of Speech Tagging using Trigram Models

$$P(t_0, \dots, t_n \mid w_0, \dots, w_n) = \prod_{i=0}^n P(w_i \mid t_i) \times P(t_i \mid t_{i-2}, t_{i-1})$$

- Each state in the *hmm* is of the form  $\langle t_j, t_k \rangle$  where  $i, j$  vary over all tags. Number of states is  $|T|^2$  for a tag set  $T$ .
- Each transition from  $\langle t_{i-2}, t_{i-1} \rangle$  to  $\langle t_{i-1}, t_i \rangle$  occurs with transition probability  $P(t_i \mid t_{i-2}, t_{i-1})$
- Each state  $\langle t_{i-1}, t_i \rangle$  emits word  $w_i$  with emission probability  $P(w_i \mid t_i)$

## Part of Speech Tagging using Trigram Models

- So, all we need to do to find the most likely tag sequence is to *train* the following two probability models:

$$P(w_i | t_i) \text{ and } P(t_i | t_{i-2}, t_{i-1})$$

- Easy to do if we have **training data** with word and tag sequences.
- All we need after we have the probability models is an algorithm to find the most likely tag sequence
- Use the algorithm used to find the best tag sequence in Hidden Markov Models: the *Viterbi* algorithm

## Part of Speech Tagging using Trigram Models

- **Evaluation:** *train* your model on the training data, *test* on unseen test data to obtain best tag sequence for each word sequence.
- **Accuracy** is measured as the percentage of correct tags for words in the test data.

# Brief History of Part of Speech Tagging

- Corpus building: English
  - Brown Corpus: 1979 (87 tags)
  - Penn Treebank Corpus: 1993 (45 tags)
  - British National Corpus (BNC): 1997
  - LOB corpus
- Other languages: Chinese, Czech, German, Korean, Turkish, . . .

# Brief History of Part of Speech Tagging

- Models and Algorithms:
  - ngram models for tagging: Church 1988
  - extension of ngram model using HMMs: Xerox (Cutting et al) 1992
  - Transformation-Based Learning: Brill 1995
  - Maximum Entropy Models: Ratnaparkhi 1997

# Applications of Part of Speech Tagging

- Other applications in NLP can be represented as POS tagging:
  - Chunking
  - Named-entity recognition (name-finding)
  - Cascaded Chunking
  - Word segmentation



# Standard Part of Speech Tagging

- Part of speech tagging: finding the best sequence of POS tags for an input sentence (word sequence)
  - Representation: what does each POS tag represent?
  - Tagset: standard POS tags (NN=noun, VB=verb, etc.)
  - Training: word sequences with corresponding tag sequences
  - Input: word sequences (sentence)
  - Output: tag sequence

# Noun Phrase Chunking

- Noun phrase chunking: e.g. input: *The man the news demonized . . .*,  
output: [ The man ] [ the news ] demonized . . .
  - Representation: is each word inside an NP or not?
  - Tagset: 3 tags: **I** (inside NP), **O** (outside NP), **B** (boundary of 2 NPs) e.g. *The/I man/I the/B news/I demonized/O . . .*
  - Training: word sequences with chunk tag sequences
  - Input: word sequences (sentence)
  - Output: chunk sequence

# Noun Phrase Chunking

- Noun phrase chunking: *The***A** *man***A** *the***B** *news***A** *demonized***O** . . .
  - Tagset: Different options for the tags, as long as they correspond to the bracketing: *[ The man ] [ the news ] demonized . . .*
  - For example, another representation could be: **I** (inside NP), **O** (outside NP), **E** (end of NP)  
e.g. *The***A** *man***E** *the***A** *news***E** *demonized***O** . . .
  - If training data is in one representation, then we can transform from one tagset to another
- What about other kinds of phrases?

# General Chunking

- Intuition for Noun Phrase chunking: In the sentence

*The company with the highest gain yesterday collapsed in today's market*

The relationship between the verb *collapsed* is to the entire phrase  
*The company with the highest gain yesterday*

- Similar intuition about other phrases, like prepositional phrases: *in today's market*

# General Chunking

- General chunking is non-overlapping:  
e.g. input: *The company with the highest gain yesterday collapsed in today's market,*  
output: [**B-NP** The company] [**B-PP** with] [**B-NP** the highest gain] [**B-NP** yesterday] [**B-VP** collapsed] [**B-PP** in] [**B-NP** today's market]
  - Representation: is each word inside a chunk or not?
  - Tagset: **O** tag for outside chunk, **B-** or **E-** prefix to the types of chunks we want, for instance **NP**, **VP**, **PP**  
e.g. *The/B-NP company/E-NP with/B-PP the/B-NP highest/B-NP gain/E-NP yesterday/B-NP collapsed/B-VP in/B-PP today's/B-NP market/B-NP*

# General Chunking

- General chunking is non-overlapping
  - Representation: is each word inside a chunk or not?
  - Tagset: **O** tag for outside chunk, **B-** or **E-** prefix to the types of chunks we want, **NP**, **VP**, **PP**
  - Training: word sequences with corresponding chunk tag sequences
  - Input: word sequences (sentence)
  - Output: chunk sequence

# Named Entity Recognition

- In the sentence

*Mr. Vinken is chairman of Elsevier N. V. , a publishing group based in the Netherlands .*

- We want to find names, such as person names, corporation names of locations:

[PER Mr. Vinken] is chairman of [ORG Elsevier N. V.] , a publishing group based in the [LOC Netherlands] .

# Named Entity Recognition

- A *named entity* is a chunk that contains only names of persons, organizations or locations
  - Representation: a word or group of words as a named entity
  - Tagset: **O** tag for outside any named entity, **B-** or **E-** prefix to the types of named entities we want: **PER** = **person**, **LOC** = **location**, **ORG** = **organization**
  - Training: word sequences with corresponding named-entity tag sequences
  - Input: word sequences (sentence)  
Output: named-entity tag sequence



## Cascaded Chunking

<b>Input:</b>	Mr.	Vinken	is	chairman	of	Elsevier	N.	V.
<b>POS:</b>	NNP	NNP	VBZ	NN	IN	NNP	NNP	NNP
<b>NP:</b>	I-NP	E-NP		I-NP		I-NP	I-NP	I-NP
<b>PP:</b>					I-PP	I-PP	I-PP	I-PP
<b>VP:</b>			I-VP	I-VP	I-VP	I-VP	I-VP	I-VP
<b>S:</b>	I-S	I-S	I-S	I-S	I-S	I-S	I-S	I-S

# Cascaded Chunking

- A sequence of tagging steps
- Each step adds some more information
- Chunking had the disadvantage of not having overlapping chunks, cascaded chunking does not have this problem  
However, later steps cannot fix errors in earlier steps. For instance, a part of speech tagging error can cause errors in every successive step of cascaded chunking
- Later we will look at trees which generalize cascaded chunking in a principled way.

# Summary: Part of Speech (POS) Tagging

- POS tagging is very similar to Hidden Markov Models
- POS tagging models are different from HMMs in the following ways:
  - The state sequences correspond to a particular representation
  - The training data always has to contain the right tag for each word in the word (or observation) sequence
- Viterbi algorithm provides the best sequence of tags for a given input  
Part of speech tagging can be applied to many applications like chunking, name finding, among others